

6.824 2020 Lecture 20: Blockstack

why are we looking at Blockstack?

it touches on three questions that interest me:

- how to build a naming system -- a PKI -- a critical missing piece

- a non-crypto-currency use of a blockchain

- a very different architecture for web sites, might someday be better

Blockstack is a real system, developed by a company

it does have some users and some apps written for it

but view it as an exploration of how things could be different and better

- not as "this is definitely how things should be"

what's a decentralized app?

apps built in a way that moves ownership of data into users's hands

- and out of centrally-controlled web sites

there are many recent (and older) explorations of this general vision.

the success (and properties) of Bitcoin has prompted a lot of recent activity

old: a typical (centralized) web site

[user browsers, net, site's web servers w/ app code, site's DB]

users' data hidden behind proprietary app code

e.g. blog posts, gmail, piazza, reddit comments, photo sharing,

- calendar, medical records, &c

this arrangement has been very successful

- it's easy to program

why is this not ideal?

- users have to use this web site's UI if they want to see their data

- web site sets (and changes!) the rules for who gets access

- web site may snoop, sell information to advertisers

- web site's employees may snoop for personal reasons

- disappointing since it's often the user's own data!

a design view of the problem:

- the big interface division is between users and app+data

- app+data integration is convenient for web site owner

- but HTML as an interface is UI-oriented.

- and is usually not good about giving users control and access to data

new: decentralized apps

[user apps, general-purpose cloud storage service, naming/PKI]

this architecture separates app code from user data

- the big interface division is between user+app and data

- so there's a clearer notion of a user's data, owned/controlled by user

- much as you own the data on your laptop, or in your Athena account

requirements for the storage system

- in the cloud, so can be accessed from any device

- general-purpose, like a file system

- paid for and controlled by user who owns the data

- sharing between users, modulo permissions, for multi-user apps

- sharing between a user's apps, modulo permissions

- similar to existing services like Amazon S3

what's the point?

- easier for users to switch apps, since data not tied to apps (or web sites)

- easier to have apps that look at multiple kinds of data

- calendar/email, or backup, or file browser

privacy vs snooping (assuming end-to-end encryption)

how might decentralized applications work?

here's one simple possibility.

app: a to-do list shared by two users

[UI x2, check-box list, "add" button]

both contribute items to be done

both can mark an item as finished

a public storage system,

key/value data owned by each of U1 and U2

users U1 and U2 run apps on their computers

maybe as JavaScript in browsers

the apps read other user's public data, write own user's data

the app doesn't have any associated server, it just uses the storage system

each user creates a file with to-do items

and a file with "done" marks

each user's UI code periodically scans the other user's to-do files

the point:

the service is storage, independent of any application.

so users can switch apps, write their own, add encryption to

prevent snooping, delete their to-do lists, back them up,

integrate with e-mail app, &c

what could go wrong?

decentralization is painful:

per-user FS-like storage much less flexible than dedicated SQL DB

no trusted server to e.g. look at auction bids w/o revealing

cryptographic privacy/authentication makes everything else harder

awkward for users as well as programmers

current web site architecture works very well

easy to program

central control over software+data makes changes (and debugging) easy

good solutions for performance, reliability

easy to impose application-specific security

successful revenue model (ads)

now for Blockstack

why does Blockstack focus on naming?

names correspond to human users, e.g. "robertmorris"

name -> location (in Gaia) of user's data, so multiple users can interact

name -> public key, for end-to-end data security

so I can check I've really retrieved your authentic data

so I can encrypt my data so only you can decrypt it

since storage system is not trusted

lack of a good global PKI has been damaging to many otherwise good security ideas

so Blockstack started with names

Blockstack claims naming is hard, summarized by "Zooko's triangle":

1. unique (global) i.e. each name has the same meaning to everyone

2. human-readable

3. decentralized

claim: all three would be valuable (debatable...)

claim: any two is easy; all three is hard

example for each pair of properties?

unique + human-readable : e-mail addresses

unique + decentralized : randomly chosen public keys

human-readable + decentralized : my contact list

why is all three hard?

can we add the missing property to any of our three schemes?

no, all seem to be immediate dead ends

summary of how Blockstack gets all three?

Bitcoin produces an ordered chain of blocks

Blockstack embeds name-claiming records in Bitcoin blocks

if my record claiming "rtm" is first in Bitcoin chain, I own it

unique (== globally the same)?

human-readable?

decentralized?

is this kind of name space good for decentralized apps?

is unique (== global) valuable?

yes: I may be able to remember names I already know.

yes: I can give you a name, and you can use it.

yes: I can look at an ACL and guess what it means.

no: human-readable names aren't likely to be very meaningful if chosen from global pool

e.g. robert_morris_1779 -- is that me? or someone else?

how about "rtm@mit.edu"?

no: how can I find your Blockname name?

how can I verify that a Blockstack name is really you?

other (possibly bad) ideas:

only public keys, don't bother with human-readable names

each person keeps separate "contact list" with names they understand

naturally decentralized

not "unique" thus no need for Bitcoin

central entity that reliably verifies human identity

what are all the pieces in Blockstack?

client, browser, application, blockstack.js

Blockstack Browser (meant to run on client machine)

Bitcoin's block-chain

Blockstack servers

read Bitcoin chain

interpret Blockstack naming records to update DB

serve naming RPCs from clients

name -> pub key + zone hash

Atlas servers -- store "zone records"

a name record in bitcoin maps to a zone record in Atlas

zone record indicates where my Gaia data is stored

keyed by content-hash, so items are immutable

you can view Atlas as just reducing the size of Blockstack's Bitcoin transactions

Atlas keeps the full DB in every server

Gaia servers

separate storage area for each user (i.e. end-users)

key -> value

backed by Amazon S3, Dropbox, &c

Gaia makes them all look the same

most users use Gaia storage provided by Blockstack

user's profile contains user's public key, per-app public keys

user can have lots of other files, containing app data

apps can sign and/or encrypt data in Gaia

S3, Dropbox, &c

back-ends for Gaia

NAME CREATION

how does one register a Blockstack name?

(<https://docs.blockstack.org/core/wire-format.html>)

the user does it (by running Blockstack software)

user must own some bitcoin

two bitcoin transactions: preorder, registration

preorder transaction

registration fee to "burn" address

hash(name)

registration transaction

name (not hashed)

owner public key

hash(zonefile)

Blockstack info hidden inside the transactions, Bitcoin doesn't look at it

but Bitcoin signatures/hashes cover this Blockstack info

why *two* transactions?

front-running

why the registration fee? after all there's no real cost.

what if a client tries to register a name that's already taken?

what if two clients try to register same name at same time?

is it possible for an attacker to change a name->key binding?

after all, anyone can submit any bitcoin transaction they like

is it possible for Blockstack to change a name->key binding?

STORAGE

how does the client know where to fetch data from?

starting with owning user's name, and a key

apps probably use well known keys, e.g. "profile" or "todo-list"

bitcoin/blockstack, hash(zone), gaia address

how does the client check that it got the right data back from Gaia?

how does the client know data from Gaia is fresh (the latest version)?

owner signed the data when writing

where can others get the owner's public key, to check signature?

how does Gaia know whether to let a client write/change/delete?

what about encryption for privacy?

if only the owner should see the data?

if one other user should see the data, in addition to the owner?

if just 6.824 students should see the data?

PRIVATE KEYS

never leaves user's device(s)

so you don't have to trust anything other than your device and Blockstack's software

each of your devices has a copy of your master private key

"master" private key only seen by Blockstack Browser

too sensitive to let apps see or use it

protected by pass-phrase, then in clear while user is active

Blockstack Browser hands out per-app private keys
so each app has more or less separate encrypted storage
makes it hard for one user's different apps to cooperate
sometimes that's what you want
sometimes you do want sharing among your own apps

DISCUSSION

here are some questions to chew on.
about naming
about decentralized applications
you can view them as criticism.
or as areas for further development.

Q: could blockstack be used as a PKI for e-mail, to map rtm@mit.edu to my public key?
blockstack names vs e-mail addresses?
what does a blockstack name mean?

Q: why is PKI hard in general?
lost pass-phrases and keys
recovery (mother's maiden name? SMS? e-mail?)
what does a name mean? connection to "real" identity?
how to go from intuitive notion of who I want to talk to, to name?
some progress, e.g. Keybase

Q: for naming and PKI, is there strong value in decentralization?
can we have a centralized but secure naming system?
who can we all trust for a global-scale system?
indeed what value can a central authority realistically deliver?
would adoption be easier with decentralization?

Q: could blockstack use a scheme like Certificate Transparency instead of Bitcoin?
CT can't resolve conflicts, only reveal them.
different CT logs may have different order
so CT can't say which came first
it's Bitcoin mining that resolves forks and forces agreement
the fee aspect of Blockstack seems critical vs spam &c, relies on cryptocurrency
in general, open block-chains only seem to make sense w/ cryptocurrency

Q: is Blockstack convenient for programmers?
all code in client, no special servers
hard to have data that's specific to the app, vs each user
indices, vote counts, front-page rankings for Reddit or Hacker News
SQL queries
cryptographic access control, groups, revocation, &c
hard to both look at other users' secrets, and keep the secrets
e.g. for eBay
maybe only worthwhile if users are enthusiastic...

Q: is decentralized user-owned storage good for user privacy?
is it better than trusting Facebook/Google/&c web sites to keep data private?
vs other users, hackers, their own employees?
can Blockstack storage providers watch what you access?
what if app, on your computer, snoops on you?
after all, it's presumably still Facebook or whoever writing the app.

is cryptographic access control really feasible?
you still have to trust the provider to preserve your data
and to serve up the most recent version
if you trust them that much, why not trust them to keep it secret too?

Q: is decentralized user-owned storage good for user control?
do users want to switch applications a lot for the same data?
do users want to use same data in multiple applications?
does either even work in general, given different app formats?

Q: will users be willing to pay for their own Gaia storage?

CONCLUSION

what do I take away from Blockstack?

I find the overall decentralization vision attractive.
the whole thing rests on a PKI -- any progress here would be great
a general-purpose mapping from all users to their public keys would be very useful
surprising that we can have decentralized human-readable name allocation
but unclear whether decentralized human-readable names are a good idea
separating cloud data from applications sounds like a good idea
but developers will hate it (e.g. no SQL).
not clear users will know or care.
not clear whether users will want to pay for storage.
end-to-end encryption for privacy sound like a good idea
private key management is a pain, and fragile
encryption makes sharing and access control very awkward
you still have to trust vendor software; not clear it's a
huge win that it's running on your laptop rather than
vendor's server.

all that said, it would be fantastic if Blockstack or something
like it were to be successful.