

6.824 2020 Lecture 15: Spark

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing Zaharia et al., NSDI 2012

why are we looking at Spark?

- widely-used for datacenter computations
- generalizes MapReduce into dataflow
- supports iterative applications better than MapReduce
- successful research: ACM doctoral thesis award

three main topics:

- programming model
- execution strategy
- fault tolerance

let's look at page-rank

- here's SparkPageRank.scala from the Spark source repository
- like the code in Section 3.2.2, with more detail

```
1  val lines = spark.read.textFile("in").rdd
2  val links1 = lines.map{ s =>
3    val parts = s.split("\\s+")
4    (parts(0), parts(1))
5  }
6  val links2 = links1.distinct()
7  val links3 = links2.groupByKey()
8  val links4 = links3.cache()
9  var ranks = links4.mapValues(v => 1.0)
10
11  for (i <- 1 to 10) {
12    val jj = links4.join(ranks)
13    val contribs = jj.values.flatMap{
14      case (urls, rank) =>
15        urls.map(url => (url, rank / urls.size))
16    }
17    ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
18  }
19
20  val output = ranks.collect()
21  output.foreach(tup => println(s"${tup._1} has rank: ${tup._2} ."))
```

page-rank input has one line per link, extracted from a big web crawl
from-url to-url
the input is vast!

page-rank output is the "importance" of each page
based on whether other important pages point to it
really models estimated probability that someone will visit each page
user model:

- 85% chance of following a link from current page
- 15% chance of visiting a random page

page-rank algorithm

- iterative, essentially simulates multiple rounds of users clicking links
- ranks (probabilities) gradually converge
- page-rank would be awkward and slow in MapReduce

my example input -- file "in":

```
u1 u3
u1 u1
u2 u3
u2 u2
u3 u1
```

I'll run page-rank in Spark (local machine, not a cluster):

```
./bin/run-example SparkPageRank in 10
u2 has rank: 0.2610116705534049 .
u3 has rank: 0.9999999999999998 .
u1 has rank: 1.7389883294465944 .
```

apparently u1 is the most important page.

let's run some of the page-rank code in the Scala interpreter

```
./bin/spark-shell
```

```
val lines = spark.read.textFile("in").rdd
-- what is lines? does it contain the content of file "in"?
lines.collect()
-- lines yields a list of strings, one per line of input
-- if we run lines.collect() again, it re-reads file "in"
val links1 = lines.map{ s => val parts = s.split("\\s+"); (parts(0), parts(1)) }
links1.collect()
-- map, split, tuple -- acts on each line in turn
-- parses each string "x y" into tuple ( "x", "y" )
val links2 = links1.distinct()
-- distinct() sorts or hashes to bring duplicates together
val links3 = links2.groupByKey()
-- groupByKey() sorts or hashes to bring instances of each key together
val links4 = links3.cache()
-- cache() == persist in memory
var ranks = links4.mapValues(v => 1.0)

-- now for first loop iteration
val jj = links4.join(ranks)
-- the join brings each page's link list and current rank together
val contribs = jj.values.flatMap{ case (urls, rank) => urls.map(url => (url, rank / urls.size)) }
-- for each link, the "from" page's rank divided by number of its links
ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
-- sum up the links that lead to each page

-- second loop iteration
val jj2 = links4.join(ranks)
-- join() brings together equal keys; must sort or hash
val contribs2 = jj2.values.flatMap{ case (urls, rank) => urls.map(url => (url, rank / urls.size)) }
ranks = contribs2.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
-- reduceByKey() brings together equal keys

-- the loop &c just creates a lineage graph.
-- it does not do any real work.

val output = ranks.collect()
-- collect() is an action.
-- it causes the whole computation to execute!
output.foreach(tup => println(s"${tup._1} has rank: ${tup._2} ."))
```

until the final collect(), this code just creates a lineage graph
it does not process the data

what does the lineage graph look like?

Figure 3

it's a graph of transform stages -- a data-flow graph

it's a complete recipe for the computation

note that the loop added to the graph -- there is not actually a cycle

there's a **new** ranks/contribs for each loop iteration

for multi-step computation, this programming model is more convenient than MapReduce

the Scala code runs in the "driver" machine of Figure 2

the driver constructs a lineage graph

the driver compiles Java bytecodes and sends them to worker machines

the driver then manages execution and data movement

what does the execution look like?

[diagram: driver, partitioned input file, workers]

- * input in HDFS (like GFS)

- * input data files are already "partitioned" over many storage servers
first 1,000,000 lines in one partition, next lines in another, &c.

- * more partitions than machines, for load balance

- * each worker machine takes a partition, applies lineage graph in order

- * when computation on different partitions is independent ("narrow"):

 - no inter-machine communication required after first read

 - a worker applies series of transformations to input stream

this is already more efficient than MapReduce

data is forwarded directly from one transformation to the next

MR would need multiple Map+Reduces

with expensive store to GFS, then re-read, between each

what about distinct()? groupByKey()? join()? reduceByKey()?

these need to look at data from **all** partitions, not just one

because all records with a given key must be considered together

these are the paper's "wide" dependencies (as opposed to "narrow")

how are wide dependencies implemented?

[diagram]

a lot like Map intermediate output in MapReduce

the driver knows where the wide dependencies are

e.g. between the map() and the distinct() in page-rank

upstream transformation, downstream transformation

the data must be "shuffled" into new partitions

e.g. bring all of a given key together

after the upstream transformation:

- split output up by shuffle criterion (typically some key)

- arrange into buckets in memory, one per downstream partition

before the downstream transformation:

- (wait until upstream transformation completes -- driver manages this)

- each worker fetches its bucket from each upstream worker

- now the data is partitioned in a different way

wide is expensive!

- all data is moved across the network

- it's a barrier -- all workers must wait until all are done

what if data is re-used?

- e.g. links4 in our page-rank

- by default, must be re-computed, e.g. re-read from input file

- persist() and cache() cause links to be saved in memory for re-use

re-using persisted data is another big advantage over MapReduce

Spark can be optimized based on its view of the whole lineage graph

- stream records, one at a time, through sequence of narrow transformations

- increases locality, good for CPU data caches

- avoids having to store entire partition of records in memory

- notice when shuffles aren't needed b/c inputs already partitioned in the same way

- e.g. links4.join(ranks)

what about fault tolerance?

- what if one machine crashes?

- its memory and computation state are lost

- driver re-runs transformations on crashed machine's partitions on other machines

- usually each machine is responsible for many partitions

- so load can be spread

- thus re-computation is pretty fast

- for narrow dependencies, only lost partitions have to be re-executed

what about failures when there are wide dependencies?

- re-computing one failed partition requires information from *all* partitions

- so *all* partitions may need to re-execute from the start!

- even though they didn't fail

- Spark supports checkpoints to HDFS (like GFS) to cope with this

- driver only has to recompute along lineage from latest checkpoint

- for page-rank, perhaps checkpoint ranks every 10th iteration

limitations?

- geared up for batch processing of bulk data

- all records treated the same way

- transformations are "functional" -- turn input into output

- no notion of modifying data in place

summary

- Spark improves expressivity and performance vs MapReduce

- giving the framework a view of the complete dataflow is helpful

- performance optimizations

- failure recovery

- what were the keys to performance?

- leave data in memory between transformations, vs write to GFS then read

- re-use of data in memory (e.g. links in page-rank)

- Spark very successful, widely used